
Contents

1	Introduction	1
2	Related Work	1
3	Dataset description	1
4	Methods Description	1
4.1	Graph Convolutional Networks (GCN)	1
4.2	Graph Attention Networks (GAT)	2
4.3	Graph Sample and Aggregate (GraphSAGE)	2
4.4	Graph Transformer	2
5	Evaluation Strategy	3

Team members: 1. Jędrzej Miczke 156068 2. Mateusz Nowicki 156064 3. Krzysztof Skrobała 156039 4. Wojciech Bogacz 156034

1 Introduction

This project aims to test methods for predicting a paper category based on its abstract, and the graph structure of citations between papers. The main idea is to reproduce the results from the paper *Semi-Supervised Classification with Graph Convolutional Networks* by Thomas N. Kipf and Max Welling, presented at ICLR 2017, and to compare it with other methods.

2 Related Work

3 Dataset description

4 Methods Description

4.1 Graph Convolutional Networks (GCN)

Graph Convolutional Networks (GCNs) extend convolution to graph-structured data by iteratively aggregating and transforming features from a node's neighbors using the graph topology.

A common layer-wise propagation rule is:

$$H^{(l+1)} = \sigma \left(\tilde{D}^{-\frac{1}{2}} \tilde{A} \tilde{D}^{-\frac{1}{2}} H^{(l)} W^{(l)} \right)$$

where:

- $\tilde{A} = A + I$ is the adjacency matrix with self-loops.
- \tilde{D} is the degree matrix of \tilde{A} .
- $H^{(l)}$ represents the node features at layer l .
- $W^{(l)}$ denotes the learnable weights.

-
- $\sigma(\cdot)$ is a non-linear activation function.

We can note that when there are no edges, the $\tilde{D}^{-\frac{1}{2}} \tilde{A} \tilde{D}^{-\frac{1}{2}}$ term becomes the identity matrix, and the layer reduces to a standard feedforward layer, which will be useful for a baseline comparison.

4.2 Graph Attention Networks (GAT)

Graph Attention Networks (GATs) incorporate an attention mechanism to assign different importance to neighboring nodes when aggregating features.

For a node i , attention coefficients are computed as:

$$\alpha_{ij} = \frac{\exp(\text{LeakyReLU}(\vec{a}^\top [W\vec{h}_i \| W\vec{h}_j]))}{\sum_{k \in \mathcal{N}(i)} \exp(\text{LeakyReLU}(\vec{a}^\top [W\vec{h}_i \| W\vec{h}_k]))}$$

and the node update is:

$$\vec{h}'_i = \sigma \left(\sum_{j \in \mathcal{N}(i)} \alpha_{ij} W\vec{h}_j \right)$$

where: - \vec{h}_i represents the input node features. - W and \vec{a} are learnable parameters. - $\mathcal{N}(i)$ denotes the neighbors of node i . - $\|$ is the concatenation operation. - $\sigma(\cdot)$ is a non-linear activation function.

4.3 Graph Sample and Aggregate (GraphSAGE)

GraphSAGE is an inductive graph representation learning method that learns node embeddings by sampling and aggregating features from a fixed-size set of neighboring nodes.

For node i at layer l , the neighborhood aggregation is:

$$h_{\mathcal{N}(i)}^{(l)} = \text{AGGREGATE}^{(l)} \left(\{h_j^{(l)} : j \in \mathcal{N}(i)\} \right)$$

and the node update is:

$$h_i^{(l+1)} = \sigma \left(W^{(l)} \cdot [h_i^{(l)} \| h_{\mathcal{N}(i)}^{(l)}] \right)$$

followed by normalization.

where: - $\mathcal{N}(i)$ denotes the **sampled** neighbors of node i . - $W^{(l)}$ are the learnable weights. - $\|$ indicates concatenation. - $\text{AGGREGATE}(\cdot)$ is a differentiable function (e.g., mean, max-pooling, or LSTM).

4.4 Graph Transformer

Graph Transformers adapt the self-attention mechanism of Transformers to graph-structured data, enabling nodes to attend to other nodes based on learned attention scores while incorporating graph structure via positional or edge encodings.

For a node i , attention is computed as:

$$\text{Attn}(i, j) = \frac{(W_Q h_i)(W_K h_j)^\top}{\sqrt{d}}$$

and the node update is:

$$h'_i = \sum_{j \in \mathcal{V}} \text{softmax}_j (\text{Attn}(i, j) + b_{ij}) W_V h_j$$

where: * h_i are the node features. * W_Q, W_K, W_V are learnable projection matrices for Query, Key, and Value. * d is the attention dimension (used for scaling). * b_{ij} is a bias term encoding graph structure (e.g., shortest path distance or edge connectivity). * \mathcal{V} is the set of all nodes in the graph.

5 Evaluation Strategy

We measure the performance of each model using accuracy on a held-out test set. Additionally, we monitor training and validation loss to assess convergence and potential overfitting.